

CSCB09 2022 Winter – Assignment 3

In this assignment, you will practice using Unix system calls for running programs and file redirection.

Build Your Own Shell

In this assignment, you will implement a toy shell that interprets (runs) a small subset of shell commands. Overview: 3 cases: (1) echo, (2) fork-exec a program, (3) list of commands. Moreover, in each case, at each level, there may be stdout redirection.

Thus there can be an example equivalent to:

```
{
    echo A
    { ls -l ; echo B > f1 ; cat f1 ; } > f2
    echo C
}
```

So every time you redirect, you must expect to restore later, and there can be many nested levels. Commands are represented by struct cmd in byos.h. You will implement this shell with a key function -- The interpreter function. This function accepts a struct with a list of commands the shell should run. The interpreter function you will implement is

```
int interp ( const struct cmd * c)
```

Its behaviour and return value are as follows.

Command Semantics (Behaviour)

Firstly, each command may optionally specify stdout redirection to a pathname. In this case, the file should be opened for write-only, truncated if it already exists, created if not. (You are encouraged to use `open` or `creat` with octal number 0666 for the mode argument, rather than `fopen`.) If opening fails, the command should not be run, and the return value must be 1 (analogous to real shells). If opening succeeds, then perform stdout redirection, taking care to save the original stdout so you can restore later. And then, the command is one of 3 types:

- Echo: For simplicity, there is only one C string in struct `echo_d` to print to stdout; if newlines are intended, they are already in the string. You may assume that writing succeeds. You are encouraged to use `write` rather than `printf` or `puts`—since we are doing stdout redirection left right and centre, bookkeeping data of `stdio.h` functions may

become terribly invalid. The return value of this case is 0. (Optional: You may return 1 if you detect that writing fails, analogous to real shells.)

- Forx: This is the fork-exec case. `struct forx_d` has the program pathname and the command line arguments, in a format ready for straight passing to a suitable `exec` syscall (which is not `execlp`). Please also choose one that can search `PATH` so that the pathname can be simply `"ls"` for example.

You may assume that `fork` succeeds.

In the child, if `exec` fails, you may print an error message of your choice to `stderr`. Then the child must exit (why?), and the exit status must be 127 (analogous to real shells).

The parent must use `wait` to wait for the child to terminate. Then the return value must be:

- If the child exits, the return value is the child's exit status.
- If the child is killed by signal, the return value is 128+signal (analogous to real shells).
- You may ignore the other 2 cases (`STOP` and `CONT`).
- List: Multiple commands in an array to run sequentially; wait for one to finish before running the next. (If you use a recursive call, this is trivial.) The return value must be the return value from running the last command; if `n = 0`, so there is no command, the return value must be 0.

Except: If running a command results in the return value of 128+SIGINT (so the command or a descendent is killed by SIGINT), do not run the remaining commands. The return value must be 128+SIGINT. (This is analogous to real shells, `ctrl-c` stops the whole command tree.)

(So, in all cases, the return value must be that of the last run command.)

After the command is done, restore the original stdout (if redirection was done).